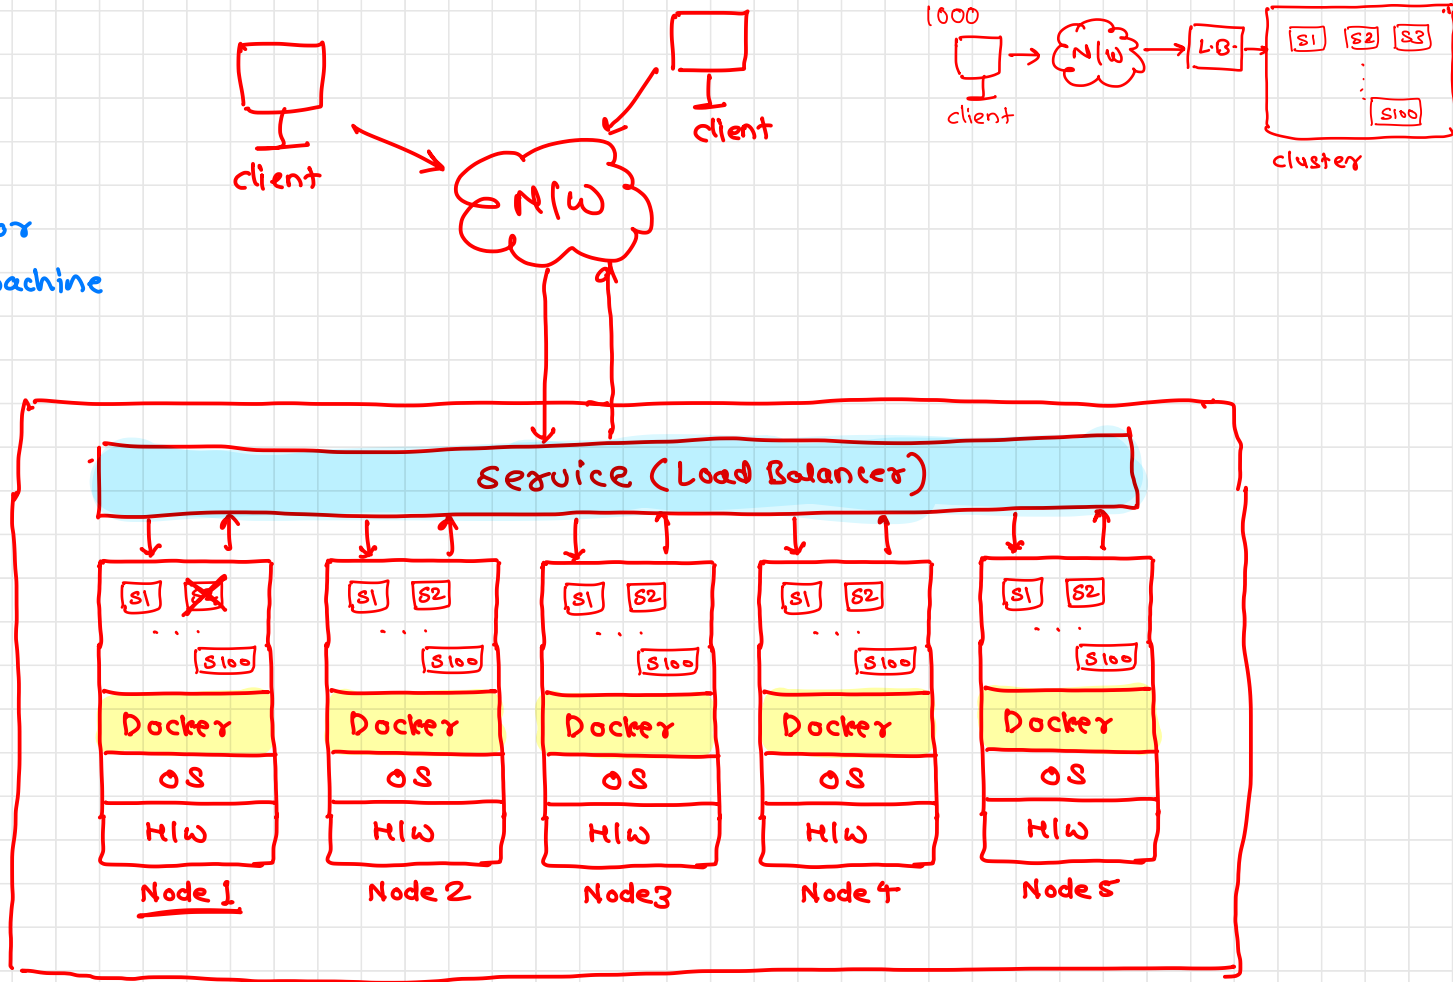


Container Orchestration



Node:

- physical or
- virtual machine



Overview

- Container orchestration is all about managing the lifecycles of containers, especially in large, dynamic environments → starting / restarting / deleting
- Software teams use container orchestration to control and automate many tasks
 - Provisioning and deployment of containers
 - Redundancy and availability of containers
 - Scaling up or removing containers to spread application load evenly across host infrastructure
 - Movement of containers from one host to another if there is a shortage of resources in a host, or if a host dies
 - Allocation of resources between containers
 - External exposure of services running in a container with the outside world
 - Load balancing of service discovery between containers
 - Health monitoring of containers and hosts
 - Configuration of an application in relation to the containers running it



Orchestration Tools

✓ ▪ Docker Swarm

✓ ▪ Kubernetes

✓ ▪ Mesos

✓ ▪ Marathon



Docker Swarm



Overview

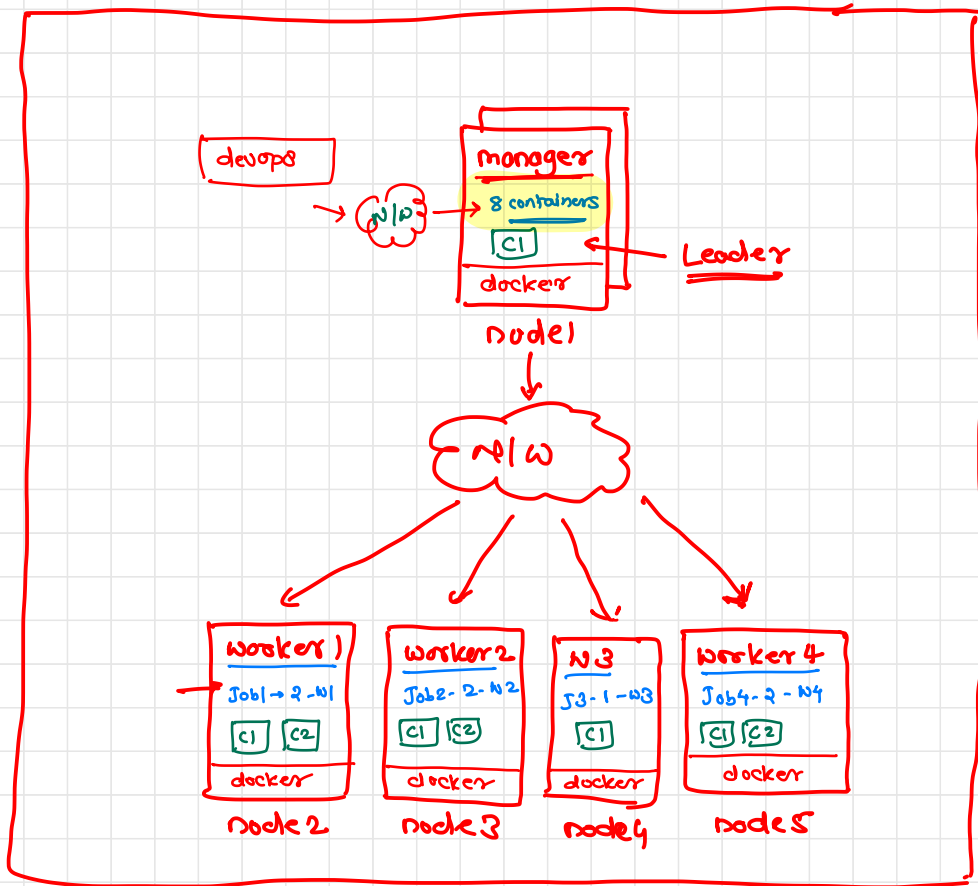
- Docker Swarm is a container orchestration engine
- It takes multiple Docker Engines running on different hosts and lets you use them together *cluster / docker swarm*
- The usage is simple: declare your applications as stacks of services, and let Docker handle the rest
- Services can be anything from application instances to databases
↳ *frontend / backend / database*



orchestrator



workers



SWARM

What is a swarm?

- A swarm consists of multiple Docker hosts which run in **swarm mode**
- A given Docker host can be a **manager**, a **worker**, or perform both roles
- When you create a **service**, you define its **optimal state** [no of containers]
- Docker works to maintain that desired state
 - For instance, if a worker node becomes unavailable, Docker schedules that node's tasks on other nodes
- A **task** is a running container which is part of a swarm service and managed by a swarm manager, as opposed to a standalone container
- When Docker is running in swarm mode, you can still run standalone containers on any of the Docker hosts participating in the swarm, as well as swarm services
- A key difference between standalone containers and swarm services is that only swarm **managers** can manage a swarm, while standalone containers can be started on any daemon



Features

- Cluster management integrated with Docker Engine
 - Decentralized design
 - Declarative service model
 - Scaling
 - Desired state reconciliation
 - Multi-host networking
 - Service discovery
 - Load balancing
 - Secure by default
 - Rolling updates
- [No external installation is needed]

Service - 100 containers
20x ↓
 (20) (80)



Nodes

- A **node** is an instance of the Docker engine participating in the swarm
- You can run one or more nodes on a single physical computer or cloud server
- To deploy your application to a swarm, you submit a service definition to a **manager node**
- **Manager Node**
 - The manager node dispatches units of work called tasks to worker nodes
 - Manager nodes also perform the orchestration and cluster management functions required to maintain the desired state of the swarm
 - Manager nodes elect a single leader to conduct orchestration tasks
- **Worker nodes**
 - Worker nodes receive and execute tasks dispatched from manager nodes
 - An agent runs on each worker node and reports on the tasks assigned to it
 - The worker node notifies the manager node of the current state of its assigned tasks so that the manager can maintain the desired state of each worker



Services and tasks

■ Service

- A service is the definition of the tasks to execute on the manager or worker nodes
- It is the central structure of the swarm system and the primary root of user interaction with the swarm
- When you create a service, you specify which container image to use and which commands to execute inside running containers

■ Task

- A task carries a Docker container and the commands to run inside the container
- It is the atomic scheduling unit of swarm
- Manager nodes assign tasks to worker nodes according to the number of replicas set in the service scale
- Once a task is assigned to a node, it cannot move to another node
- It can only run on the assigned node or fail



Swarm Setup

- **Create swarm**

- > docker swarm init --advertise-addr <MANAGER-IP>

- **Get current status of swarm**

- > docker info

- **Get the list of nodes**

- > docker node ls



Swarm Setup

- **Get token (on manager node)**
> docker swarm join-token worker
- **Add node (on worker node)**
> docker swarm join --token <token>



Swarm Service

- **Deploy a service**

> docker service create --replicas <no> --name <name> -p <ports> <image> <command>

- **Get running services**

> docker service ls

- **Inspect service**

> docker service inspect <service>

- **Get the nodes running service**

> docker service ps <service>



Swarm Service

- **Scale service**

> docker service scale <service>=<scale>

- **Update service**

> docker service update --image <image> <service>

- **Delete service**

> docker service rm <service>

